

# Towards Generalizable Reinforcement Learning for Trade Execution

Chuheng Zhang, Xiaoyu Chen, Yitong Duan, Jianyu Chen, Jian Li,  
Li Zhao


Tsinghua University, IIS

Microsoft Research Asian

(This is a working draft. Feel free to discuss~)

# Short Bio

- Chuheng Zhang (张楚珩) is researcher at MSRA
- Ph.D. at Institute of interdisciplinary Information Sciences, Tsinghua University (Sep, 2016 - July, 2022), advised by Prof. Jian Li (李建)
- Bachelor at Physics Department at Nanjing University
- Research: reinforcement learning applications (e.g., quantitative investment, inventory management, advertising)
- Publications: AAAI, ICLR, WWW, SIGIR, CIKM, ICDM



# Outline

- RL background
- Trade execution background
- Methods
- Results
- Discussion

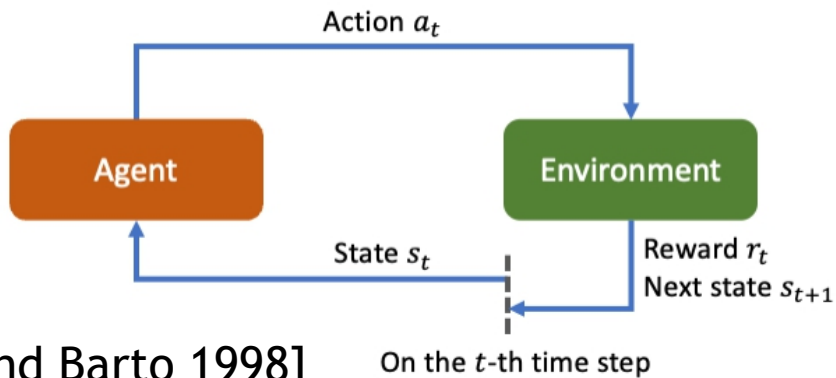
# Motivation

- RL has awesome performance on various games.
- Sequential decision is common in industry.
- Why not use RL in real applications?


Challenges	Solutions
Costly Interaction	RL from Offline Data
Sample Inefficiency	Representation Learning, Knowledge Transfer, ...
Instability	Stable Algorithm
Generalization	This work

# RL and Context

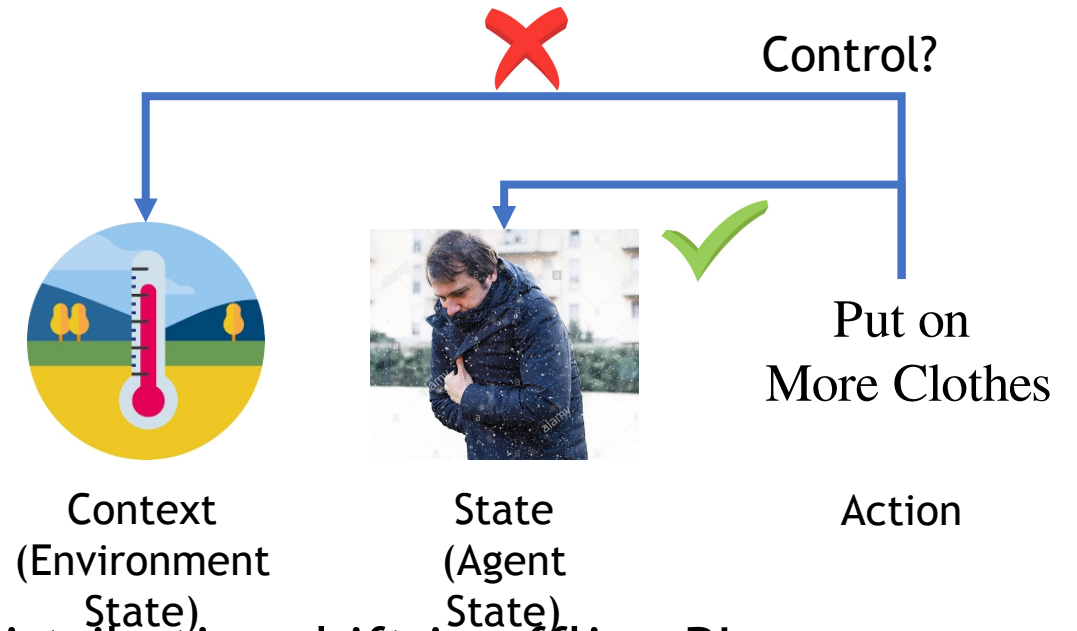
- Popular RL Applications



[Sutton and Barto 1998]



- Some Industrial RL Applications

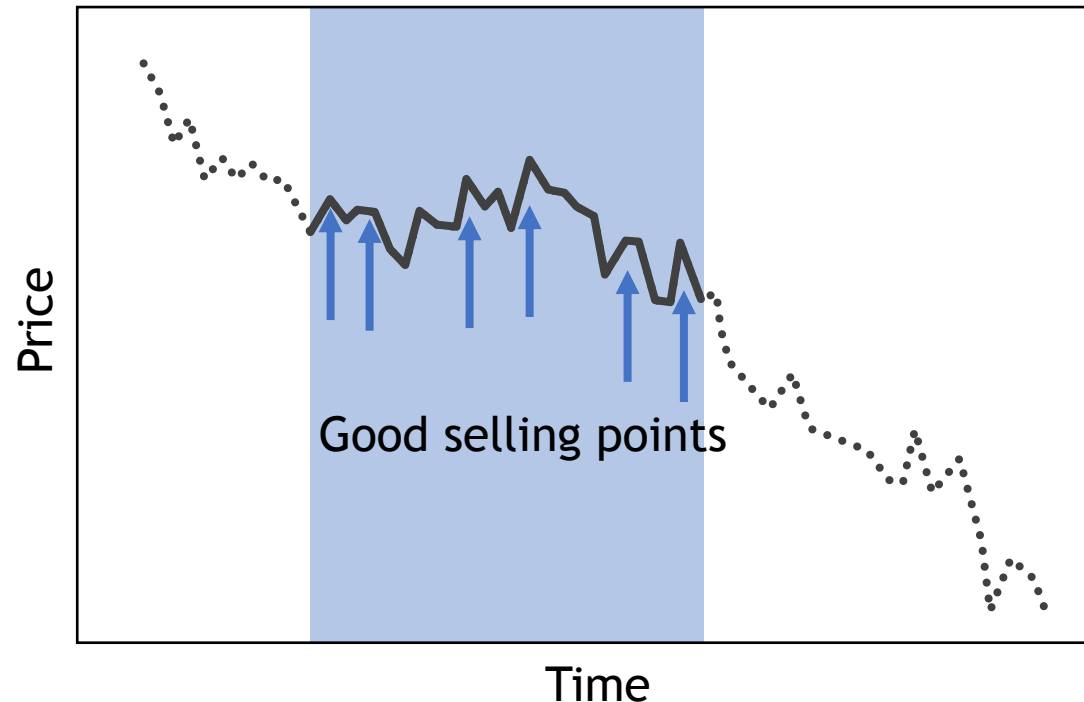


- Distribution shift in offline RL  
[Fujimoto et al 2019; Kumar et al 2020]
- Control or real-time adaptation?

# Introduction for Trade Execution

- **Trade execution** is to sell/buy a given amount of assets in a given time with the lowest possible trading cost
- **Background:** Trade execution is an **important** task for brokerage firms

Task: sell 1000 stocks within this time period



# The Setting: Limit Order Book (LOB)


- **Limit order book (LOB)** is a record of outstanding limit orders
- **Market order (MO)** is executed immediately with large trading cost
- **Limit order (LO)** ensures the price at the risk of non-execution
- **Raw offline data** contains information for us to reconstruct the LOB dynamics



Image from [Chávez-Casillas and Figueroa-López 2017]

# Background for Trade Execution


- **Traditional methods** [see e.g., Almgren and Chriss 2000; Guéant et al 2012]
  - Generate static policies
  - Rely on strong assumptions
- **Previous RL methods** suffer from overfitting
  - Observation (market indicators) is high noisy
  - Agent overfits spurious noise/feature
- **Objective:** Developing generalizable RL ag



nd



# RL formulation

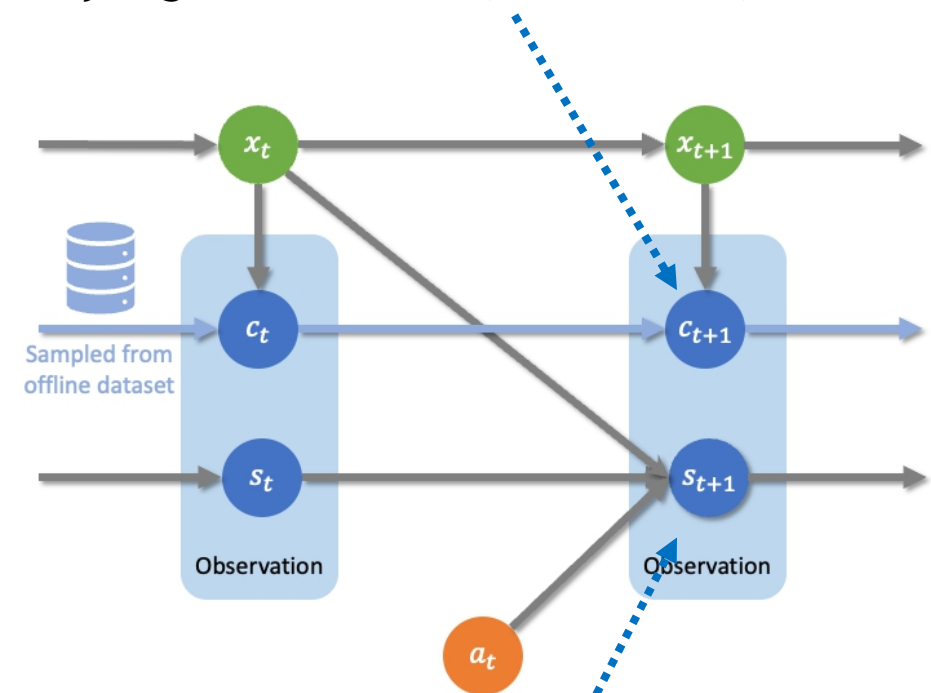
- **Simulator:** based on LOB data
- **Observation:**
  - State: inventory level, remaining time
  - Context 
- **Action:** LO (price, direction, quantity), MO (direction, quantity)
- **Reward/objective:** Maximize the cash obtained from selling the stocks
- **Challenge:** Context is high-dimensional, redundant, and noisy

	AskPrice	Volume	MACD	Spread	...
$t - k$					
...					
$t$					

# Modeling

- Offline RL with Dynamic Context (ORDC)
- $\phi: \mathcal{C} \text{ (context)} \rightarrow \mathcal{X} \text{ (latent context)}$  with  $|\mathcal{C}| \gg |\mathcal{X}|$ . E.g., for Brownian motion,  $c$  is historical prices and  $x = (\alpha, \sigma)$
- Context dynamics is highly stochastic and we assume **state dynamics is known**

**Context:** Uncontrollable  
E.g., market indicators  
Usually high-dimensional, redundant, and stochastic



**State:** Controllable  
E.g., inventory level  
The transition dynamics is usually clear

# Theoretical Analysis

- Generalization is hard when data is limited, even if the context is compressible

Theorem 1 [Sample complexity lower bound]: In a class of ORDC models, any algorithm needs at least  $\Omega\left(\frac{|\mathcal{C}|\log(|\mathcal{C}|/\delta)}{(1-\gamma)^3\epsilon^2}\right)$  context samples to learn.

- If we could properly compress the context, limited data leads to good generalization

Theorem 2 [Sample complexity upper bound]: With access to  $\phi: \mathcal{C} \rightarrow \mathcal{X}$ , an algorithm can learn with  $O\left(\frac{|\mathcal{X}|^2\log(|\mathcal{X}|/\delta)}{(1-\gamma)^4\epsilon^2}\right)$

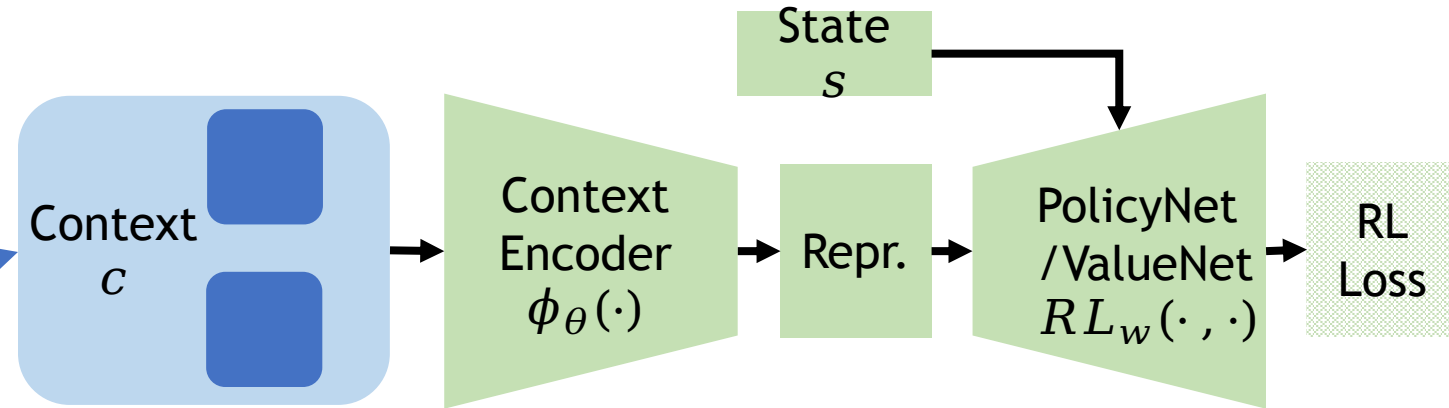
	AskPrice	Volume	MACD	RSI	...
$t-k$	Context $c \in \mathcal{C}$				
...					
$t$					

- Motivate:** Compress context representation

# Methods: RL Formulation

- Standard RL model for trade execution

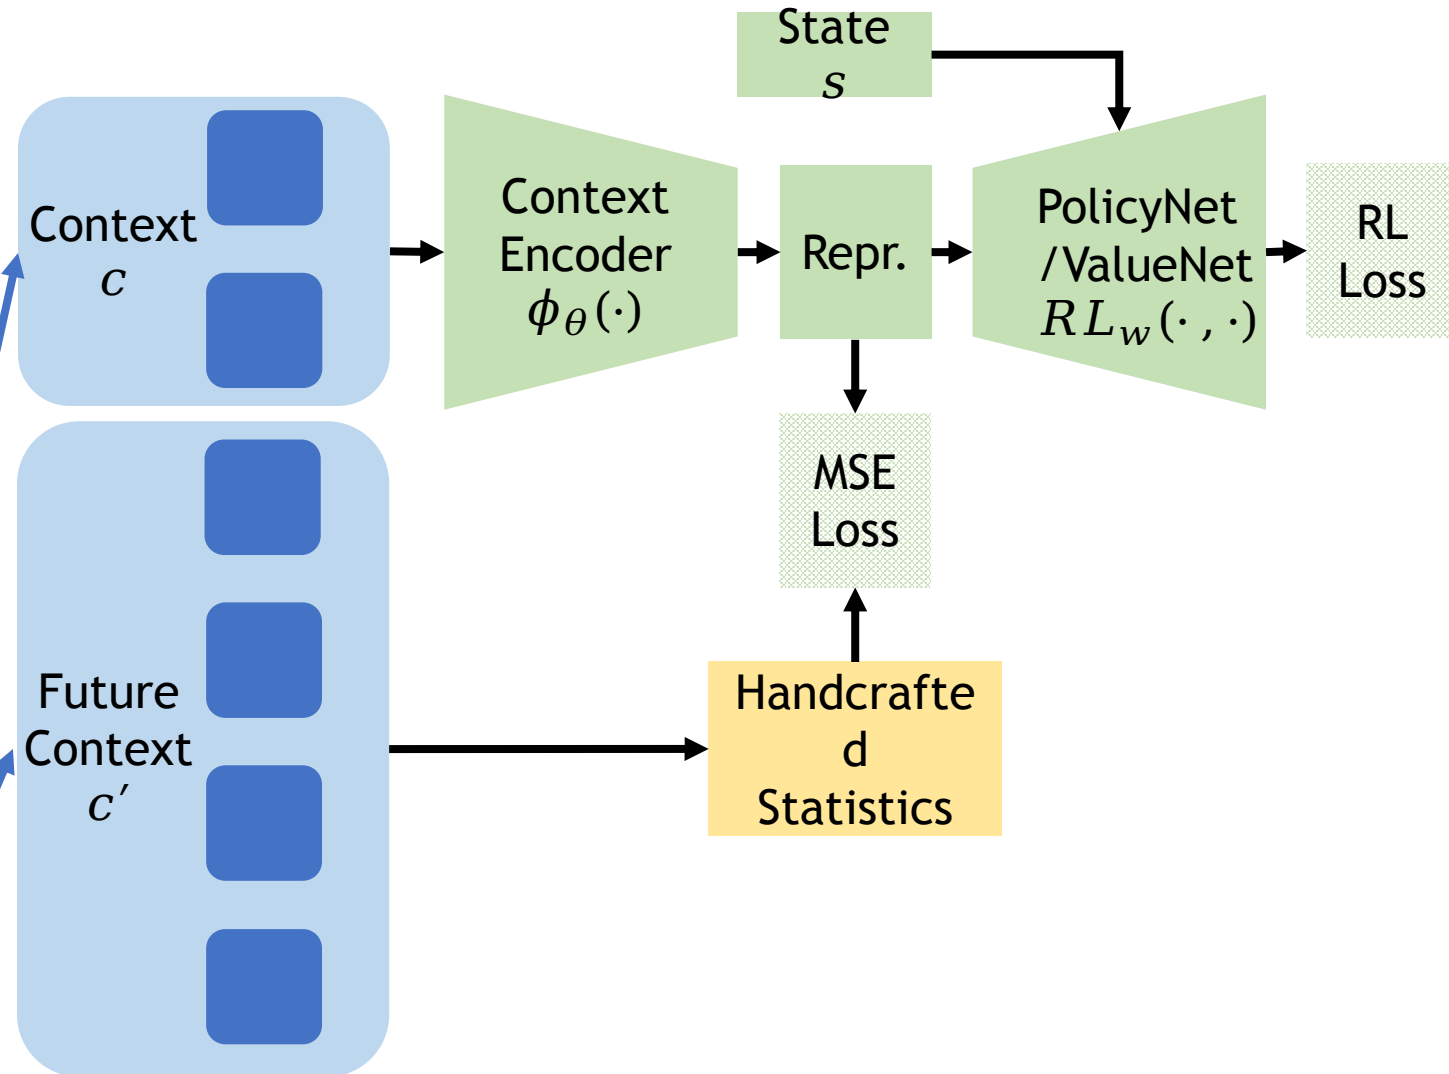
	AskPrice	Volume	MACD	...
$t - k$				
...				
$t$				



# Methods: CASH

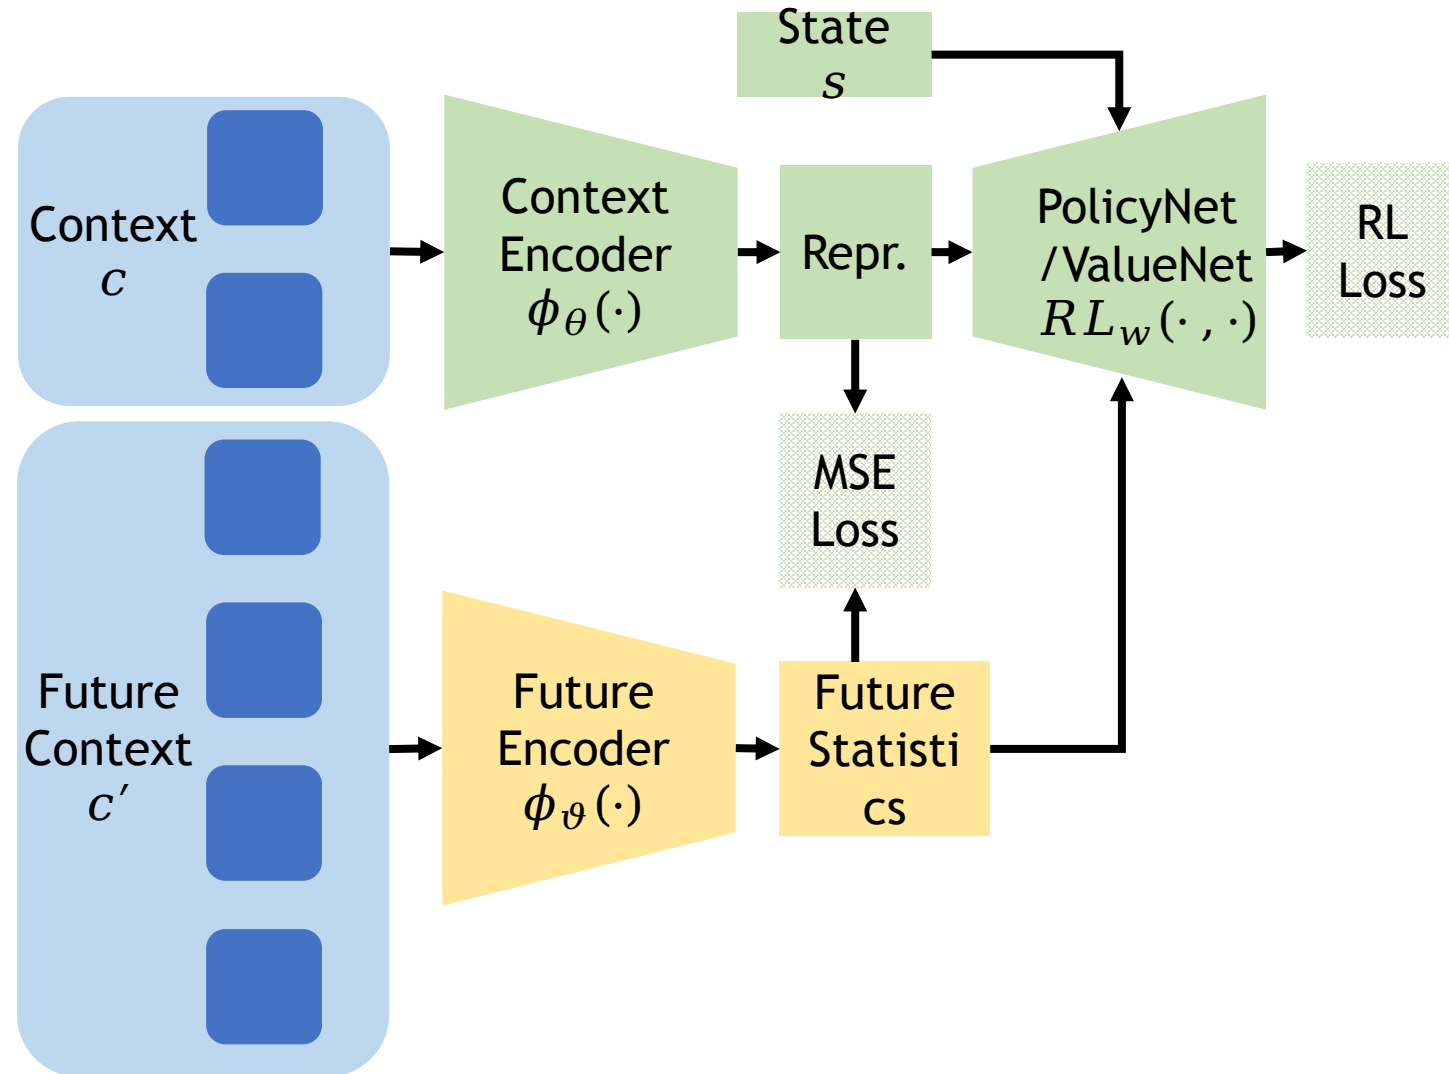
- **CASH: Context Aggregation with Handcrafted Statistics**
- Pre-train the context encoder to predict key future statistics
- E.g., future price trend and volatility

	AskPrice	Volume	MACD	...
$t - k$				
...				
$t$				
$t + 1$				
...				
$t + k'$				



# Methods: CATE

- CATE: Context Aggregation with End-to-end Training
- Future encoder is trained to
  1. Extract useful information from future
  2. Be predictable based on current context



# Methods

- Why learn a compressed context representation to predict the statistics of the future context?
  1. **Convenient:** The context representation depending on the future context is independent of actions
  2. **Informative:** Information on the future context (e.g., prices) is useful for the agent to make decisions
  3. **Feasible:** The statistics of future context is more predictable than the full future context

# Experiment Results: Trading Cost

Algorithm	Training	Validation	Testing	Generalization Gap
TWAP Strategy	-	-	13.1217 (2.0858)	-
Tuned DQN	1.6078 (2.1974)	4.9241 (2.2437)	6.7199 (2.7776)	5.1121
Nevmyvaka et al. (2006)	2.8726 (5.6576)	8.8261 (1.6331)	9.7045 (1.2583)	6.8319
Ning et al. (2018)	7.2350 (5.5857)	10.5678 (1.8090)	10.0825 (2.1145)	2.8475
Lin et al. (2020)	6.0305 (6.7465)	11.1709 (1.1422)	12.4054 (1.3553)	6.3749
Tuned DQN + CASH	1.9351 (1.8016)	3.2507 (1.2944)	3.9184 (1.2287)	<b>1.9833</b>
Tuned DQN + CATE	-2.7526 (1.3319)	-1.7554 (1.1277)	<b>0.0749</b> (1.5233)	2.8275
Tuned PPO	-1.7340 (1.7653)	1.9763 (0.5393)	3.2686 (0.5302)	5.0026
Dabérius et al. (2019)	6.7671 (11.0498)	10.2517 (2.2711)	12.9987 (2.5579)	6.2316
Lin et al. (2021)	5.5028 (7.4558)	7.8800 (1.0740)	10.2599 (1.5293)	4.4571
Fang et al. (2021)	-5.4925 (15.8669)	9.9096 (4.2206)	11.5612 (5.3556)	17.0537
Tuned PPO + CASH	-4.7082 (0.8733)	-4.1690 (0.2108)	-3.8524 (0.2234)	<b>0.8558</b>
Tuned PPO + CATE	-5.1725 (0.6705)	-4.7672 (0.1717)	<b>-4.3000</b> (0.4813)	0.8725

Table      The trading cost ( $\text{bp}=10^{-4}$ ) of different algorithms. The validation set is used for hyperparameter tuning. The numbers are the average mean (std.) trading cost in the last 10% iterations over five different random seeds.



# Experiment Results: Learned Policy

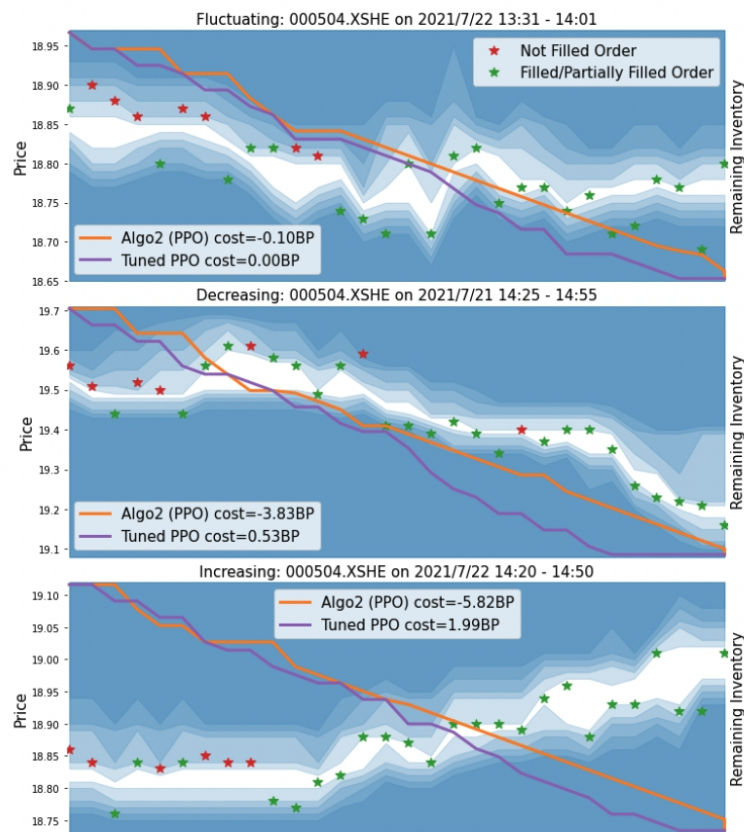


Figure 3: The policy learned by CATE (based on PPO) and the corresponding baseline. The background shaded areas indicate the 5-level ask/bid prices, and the stars indicate the orders placed by CATE. The lines indicate the remaining inventory of CATE and the baseline algorithm.

Predicted Statistics	Price	Volume
Twap volatility	1.2871	-0.4335
Avg future twap - current twap	4.5296	-0.0420
Max future twap - current twap	1.3932	-0.1761
Min future twap - current twap	-1.0587	0.2978
Sprd volatility	6.2039	-0.6810
Avg future sprd - current sprd	-0.0944	0.0879
Max future sprd - current sprd	1.8254	-0.2629
Min future sprd - current sprd	-1.2859	0.1579

Table 3: The impact of predicted statistics on the agent's action in CASH (based on PPO).

# Discussion

- Generalization in Finance
  - Data augmentation
  - Ensemble
  - Market style detection and adaptation
  - ...